

RacquetAI — Token Usage & Cost Report

Prepared by **Claude** (Anthropic Claude Code) for management, as the companion to the *Build Day Report (Made by Claude)*.
Repository: github.com/Karanvir1729/RacquetAI · Open pull request: [#1](#)
Figures as of Aug 16, 2026, ~3:05 AM EDT. The reporting session was still running when measured, so its row (and the totals) undercount slightly.

Headline numbers

302.5M TOKENS PROCESSED	2.49M TOKENS GENERATED	4 + 36 SESSIONS + SUBAGENTS	\$424–498 EST. COST, API LIST RATES
-----------------------------------	----------------------------------	---------------------------------------	---

Metric	Value
Total tokens processed (input + output + cache)	302,523,454
Tokens generated (output)	2,491,998
Fresh (uncached) input tokens	117,824
Cache writes	13,245,962
Cache reads	286,667,670
Claude sessions	4
Subagent transcripts (workflow agents)	36
Estimated cost at Anthropic API list rates	US\$424 – 498

The striking shape of these numbers — **95% of all tokens are cache reads** — is what makes a build like this economical. Claude re-reads its full working context on every step, but prompt caching serves those tokens at roughly one-tenth the normal input price.

Per-session breakdown

Tokens include each session’s own transcript **plus all subagent/workflow transcripts it spawned**. The forked session shares its first ~26 messages with the main session; those were counted once (deduplicated by message ID).

Session (models)	When (EDT)	Input	Output	Cache write	Cache read	Total
Main build — “RacquetAI repo setup and planning” (Fable 5; subagents Fable 5 + Opus 5)	Aug 15, 6:54 PM – 10:33 PM	43,854	1,009,732	4,957,760	146,124,171	152,135,517
Forked planning session — board + squash pivot (Fable 5 → Opus 5)	Aug 15, ~7:10 PM – 10:10 PM	73,393	987,382	6,091,751	113,511,819	120,664,345
Brand redraw worktree — “Redraw RacquetAI mark as a squash racquet” (Opus 5)	Aug 15, 10:06 – 10:23 PM	215	190,795	276,044	13,276,254	13,743,308
This reporting session (Fable 5 + a 4-agent mining workflow)	Aug 16, 2:38 AM –	362	304,089	1,920,407	13,755,426	15,980,284
Total		117,824	2,491,998	13,245,962	286,667,670	302,523,454

Subagents did the heavy lifting: of the ~2.49M generated tokens, roughly **1.9M were produced by the 36 workflow subagents** (feature builders, integrator, auditor, board authors, critics, icon designers, transcript miners) rather than the four top-level sessions.

Cost estimate by model

Priced at Anthropic API list rates (source: Anthropic pricing documentation): **Claude Fable 5** \$10 / \$50 per million input/output tokens; **Claude Opus 5** \$5 / \$25. Cache reads bill at ~0.1× the input rate; cache writes at 1.25× (5-minute TTL) to 2× (1-hour TTL) — the range below spans those two write assumptions.

Model	Input	Output	Cache write	Cache read	Est. cost
Claude Fable 5	35,918	1,158,659	6,420,139	131,338,310	\$270 – 318
Claude Opus 5	81,906	1,333,339	6,825,823	155,329,360	\$154 – 180
Total					\$424 – 498

Two framing notes for management:

1. **This is list-rate equivalent value, not necessarily the invoice.** Work ran through a Claude Code subscription; the figures answer “what would this compute cost at published API prices,” which is the fair basis for a build-vs-buy comparison.
2. **Counterfactual:** without prompt caching, the same 286.7M cache-read tokens would have billed as fresh input — roughly **\$2,300 of additional spend**. Caching cut the effective compute bill by ~80%.

What the tokens bought

One evening of spend (~\$424–498 list-rate) produced: a public GitHub repo with 18 commits and ~38K insertions; a working, simulator-verified iOS app with 178 passing tests; four merged feature branches plus an open PR; a complete brand system; 8 documentation files including an App Store runbook and a budget-gated AI plan; and a rendered 7-day execution board. Details and screenshots: *Build Day Report (Made by Claude)*.

Methodology

- **Source of truth:** the local Claude Code session transcripts (JSONL) for all four sessions, plus every subagent/workflow transcript they spawned (36 files, ~40 MB total). Each assistant message records exact `input_tokens`, `output_tokens`, `cache_creation_input_tokens`, and `cache_read_input_tokens` as billed by the API.
- **Deduplication:** the forked session's transcript embeds a copy of the main session's first ~26 messages; all sums are deduplicated by message UUID so no API call is counted twice.
- **Session identification:** sessions were matched to the work via timestamps against git commit times and their own recorded commands (e.g., the main session's final write coincides with commit `1f48d7a` at 10:32 PM EDT).
- **Pricing:** Anthropic list rates as of Aug 2026 (Fable 5 \$10/\$50, Opus 5 \$5/\$25 per MTok; cache read 0.1×; cache write 1.25–2×).